

EECS 598 Deep Learning

Assignment 1

Shuyang HUANG 68621288

1.

Pr.1

1. Fully-connected layer.

$$Y = W^T X + b \quad \text{which is: } y_i = \sum_j w_{ji} \cdot x_j + b_i$$

$$\frac{\partial L}{\partial w_{ji}} = \frac{\partial L}{\partial y_i} \cdot \frac{\partial y_i}{\partial w_{ji}} = \frac{\partial L}{\partial y_i} \cdot x_j, \text{ thus } \frac{\partial L}{\partial (W^T)} = \frac{\partial L}{\partial Y} \cdot X^T$$

$$\text{Thus } \frac{\partial L}{\partial W} = \left(\frac{\partial L}{\partial (W^T)} \right)^T = \left(\frac{\partial L}{\partial Y} \cdot X^T \right)^T = X \cdot \left(\frac{\partial L}{\partial Y} \right)^T$$

$$\frac{\partial L}{\partial b_i} = \frac{\partial L}{\partial y_i} \cdot \frac{\partial y_i}{\partial b_i} = \frac{\partial L}{\partial y_i} \Rightarrow \frac{\partial L}{\partial b} = \frac{\partial L}{\partial Y}$$

$$\frac{\partial L}{\partial x_j} = \sum_i \frac{\partial L}{\partial y_i} \cdot \frac{\partial y_i}{\partial x_j} = \sum_i \frac{\partial L}{\partial y_i} \cdot w_{ji} = \sum_i w_{ji} \frac{\partial L}{\partial y_i}$$

$$\text{Thus } \frac{\partial L}{\partial X} = W \cdot \frac{\partial L}{\partial Y}$$

$$\text{Thus: } \frac{\partial L}{\partial W} = X \cdot \left(\frac{\partial L}{\partial Y} \right)^T, \frac{\partial L}{\partial b} = \frac{\partial L}{\partial Y}, \frac{\partial L}{\partial X} = W \cdot \frac{\partial L}{\partial Y}$$

2. ReLU

Notice that $Y = \text{ReLU}(X) = \max\{0, X\}$.

$$\text{Thus } y_{ij} = \begin{cases} x_{ij} & x_{ij} \geq 0 \\ 0 & x_{ij} < 0 \end{cases} \Leftrightarrow y_{ij} = x_{ij} \cdot \mathbb{1}\{x_{ij} \geq 0\}$$

$$\frac{\partial y_{ij}}{\partial x_{ij}} = \mathbb{1}\{x_{ij} \geq 0\} \Rightarrow \frac{\partial L}{\partial x_{ij}} = \frac{\partial L}{\partial y_{ij}} \cdot \frac{\partial y_{ij}}{\partial x_{ij}} = \frac{\partial L}{\partial y_{ij}} \cdot \mathbb{1}\{x_{ij} \geq 0\}$$

$$\text{Thus } \frac{\partial L}{\partial X} = \mathbb{1}\{X \geq 0\} \cdot \frac{\partial L}{\partial Y}$$

3. Dropout

$Y = X \odot M$ Notice $M \in \mathbb{B}^{m \times n}$ $\mathbb{B} = \{0, 1\}$, binary set.

$$\text{Thus } y_{ij} = x_{ij} \cdot m_{ij} \Rightarrow \frac{\partial y_{ij}}{\partial x_{ij}} = m_{ij}$$

$$\frac{\partial L}{\partial x_{ij}} = \frac{\partial L}{\partial y_{ij}} \cdot \frac{\partial y_{ij}}{\partial x_{ij}} = \frac{\partial L}{\partial y_{ij}} m_{ij}$$

$$\text{Thus } \frac{\partial L}{\partial X} = \frac{\partial L}{\partial Y} \odot M$$

4. Batch Normalization

$$y_i = \gamma \left(\frac{x_i - \mu}{\sigma} \right) + \beta, \quad \mu = \frac{1}{n} \sum_j x_j, \quad \sigma = \sqrt{\frac{1}{n} \sum_j (x_j - \mu)^2 + \epsilon}$$

$$\text{Obviously, } \frac{\partial \mu}{\partial x_i} = \frac{1}{n}$$

$$\begin{aligned} \sigma &= \sqrt{\left[\left(1 - \frac{1}{n}\right) x_i - \frac{1}{n} \sum_{j \neq i} x_j \right]^2 + \epsilon} \quad \frac{\partial \sigma}{\partial x_i} = \frac{1}{2\sigma} \cdot 2 \left[\left(1 - \frac{1}{n}\right) x_i - \frac{1}{n} \sum_{j \neq i} x_j \right] \\ &= \frac{1}{\sigma n} [x_i - \mu] \end{aligned}$$

$$\begin{aligned} \frac{\partial L}{\partial x_i} &= \frac{\partial L}{\partial y_i} \cdot \frac{\partial y_i}{\partial x_i} + \sum_j \frac{\partial L}{\partial y_j} \frac{\partial y_j}{\partial \mu} \cdot \frac{\partial \mu}{\partial x_i} + \frac{\partial y_i}{\partial \sigma} \cdot \frac{\partial \sigma}{\partial x_i} \\ &= \frac{\partial L}{\partial y_i} \cdot \frac{\gamma}{\sigma} + \sum_j \frac{\partial L}{\partial y_j} \left(-\frac{\gamma}{\sigma} \cdot \frac{1}{n} + \frac{1}{\sigma n} [x_i - \mu] \cdot \left[-\frac{1}{\sigma^3} \gamma (x_j - \mu) \right] \right) \\ &= \frac{\partial L}{\partial y_i} \cdot \frac{\gamma}{\sigma} - \frac{\gamma}{\sigma n} \sum_j \frac{\partial L}{\partial y_j} - \frac{1}{n \sigma^3} \gamma [x_i - \mu] \sum_j \frac{\partial L}{\partial y_j} (x_j - \mu) \end{aligned}$$

$$\text{Thus } \frac{\partial L}{\partial x_i} = \frac{\partial L}{\partial y_i} \cdot \frac{\gamma}{\sigma} - \frac{\gamma}{\sigma n} \sum_j \frac{\partial L}{\partial y_j} - \frac{\gamma}{n \sigma^3} [x_i - \mu] \sum_j \frac{\partial L}{\partial y_j} (x_j - \mu)$$

5. Convolution.

$$\begin{aligned}
 \left(\frac{\partial L}{\partial x_{n,c}} \right)_{ij} &= \left[\sum_{f,p,q} \frac{\partial L}{\partial y_{n,f,p,q}} \cdot \frac{\partial y_{n,f,p,q}}{\partial x_{n,c,i,j}} \right] \\
 &= \left[\sum_{f,p,q} \frac{\partial L}{\partial y_{n,f,i-p+1,j-q+1}} \cdot \frac{\partial y_{n,f,i-p+1,j-q+1}}{\partial x_{n,c,i,j}} \right] \\
 &= \left[\sum_{f,p,q} \frac{\partial L}{\partial y_{n,f,i-p+1,j-q+1}} \cdot w_{f,c,p,q} \right] \\
 &= \sum_f w_{f,c} *_{full} \frac{\partial L}{\partial y_{n,f}}
 \end{aligned}$$

$$\begin{aligned}
 \left(\frac{\partial L}{\partial w_{f,c}} \right)_{ij} &= \left[\sum_n \frac{\partial L}{\partial w_{f,c,i,j}} \right] = \left[\sum_{n,p,q} \frac{\partial L}{\partial y_{n,f,p,q}} \cdot \frac{\partial y_{n,f,p,q}}{\partial w_{f,c,i,j}} \right] \\
 &= \left[\sum_{n,p,q} \frac{\partial L}{\partial y_{n,f,i-p+1,j-q+1}} \cdot \frac{\partial y_{n,f,i-p+1,j-q+1}}{\partial w_{f,c,i,j}} \right] \\
 &= \left[\sum_{n,p,q} \frac{\partial L}{\partial y_{n,f,i-p+1,j-q+1}} \cdot x_{n,c,i,j} \right] \\
 &= \sum_n x_{n,c} *_{full} \left(\frac{\partial L}{\partial y_{n,f}} \right)
 \end{aligned}$$

3

2. Logistic Classifier

a.

```

model = LogisticClassifier(input_dim=20, hidden_dim = None, reg = 0,
weight_scale=1e-1)
controller = Solver(model, data,
                    update_rule='sgd_momentum',
                    optim_config={
                        'learning_rate': 1e0,
                    },
                    lr_decay=0.95,
                    num_epochs=160, batch_size=50,
                    print_every=100)

```

This is the parameters used, and reached an accuracy of 92% with test data.

b.

```

model = LogisticClassifier(input_dim=20, hidden_dim = 120, reg = 0,
weight_scale=1e-1)
controller = Solver(model, data,
                    update_rule='sgd_momentum',
                    optim_config={
                        'learning_rate': 1e0,
                    },
                    lr_decay=0.95,
                    num_epochs=160, batch_size=50,
                    print_every=100)

```

This is the parameters used, and reached an accuracy of 91.4% with test data.

3. SVM Classifier

a.

```

model = SVM(input_dim=20, hidden_dim = None, reg = 0, weight_scale=1e-2)
controller = Solver(model, data,
                    update_rule='sgd_momentum',
                    optim_config={
                        'learning_rate': 1e0,
                    },
                    lr_decay=0.98,
                    num_epochs=200, batch_size=50,
                    print_every=100)

```

This is the parameters used, and reached an accuracy of 92.8% with test data.

b.

```

model = SVM(input_dim=20, hidden_dim = 120, reg = 0, weight_scale=1e-2)
controller = Solver(model, data,
                    update_rule='sgd_momentum',
                    optim_config={
                        'learning_rate': 1e0,
                    },
                    lr_decay=0.95,
                    num_epochs=200, batch_size=50,
                    print_every=100)

```

This is the parameters used, and reached an accuracy of 93.2% with test data.

4. Softmax Regression

a.

```

model = SoftmaxClassifier(input_dim=28*28, hidden_dim = None, reg = 0,
num_classes=10, weight_scale=1e-3)
controller = Solver(model, data,
                    update_rule='sgd_momentum',
                    optim_config={
                        'learning_rate': 1e-5,
                    },
                    lr_decay=0.95,
                    num_epochs=10, batch_size=50,
                    print_every=100)

```

This is the parameters used, and reached an accuracy of 90.2% with test data.

b.

```

model = SoftmaxClassifier(input_dim=28*28, hidden_dim = 600, reg = 0,
num_classes=10, weight_scale=1e-3)
controller = Solver(model, data,
                    update_rule='sgd_momentum',
                    optim_config={
                        'learning_rate': 1e-3,
                    },
                    lr_decay=0.8,
                    num_epochs=4, batch_size=50,
                    print_every=100)

```

This is the parameters used, and reached an accuracy of 96.71% with test data.

5. Convolutional Neural Network

a.

```
model = ConvNet(input_dim=(1, 28, 28), hidden_dim = 600, reg = 0,
num_classes=10, weight_scale=1e-3, drop_out = False)
controller = Solver(model, data,
                    update_rule='sgd_momentum',
                    optim_config={
                        'learning_rate': 1e-3,
                    },
                    lr_decay=0.9,
                    num_epochs=15, batch_size=50,
                    print_every=1)
```

This is the parameters used, and reached an accuracy of 98.7% with test data.

b.

```
model = ConvNet(input_dim=(1, 28, 28), hidden_dim = 600, reg = 0,
num_classes=10, weight_scale=1e-3, drop_out = True)
controller = Solver(model, data,
                    update_rule='sgd_momentum',
                    optim_config={
                        'learning_rate': 1e-3,
                    },
                    lr_decay=0.9,
                    num_epochs=15, batch_size=50,
                    print_every=1)
```

This is the parameters used, and reached an accuracy of 98.9% with test data.

6. VGG11

With nothing changed, the result shows:

```
Accuracy of the network on the 10000 test images: 70 %
```

7. Short answer question

a.

Sigmoid function, $f(x) = \frac{1}{1+\exp^{-x}}$, maps the the real number range into the range $[0, 1]$. With this realized, values will more likely to be hushed onto values approaching 0 and 1, and thus saturate at the these two points. With improper initial values or improper model, sigmoid have the chance to perform pretty bad.

b.

Suppose that, at a probability p , we randomly drop a neuron, which means, at probability p , the contribution to a output will not counted into our current estimation \hat{y} . Thus, to ensure the difference, or the grandient not mismatch too much, we should consider the $(1 - p)y$ as the standard values, to reduce the mismatch.

c.

Reduce the decay, thus to ensure the learning step reduce quickly, to avoid the over-fitting.